

Task Discussion

Carsharing

Daniel Graf

ETH Zürich

December 16, 2015

This task is based on the paper "*Scheduling Transfers of Resources over Time: Towards Car-Sharing with Flexible Drop-Offs*" by Kateřina Böhmová, Yann Disser, Matúš Mihalák and Rastislav Šrámek.

Problem Statement


■ Given:

- A list of S rental stations with the number of initially available cars.
- A list of N car rental requests each with start and end time and location and its profit.

■ Wanted: Find a feasible subset of the requests that maximizes the total profit.

■ Subtasks:

- 1 Only a single car
- 2 At most 20 requests
- 3 At most 10000 requests
- 4 Many different times
- 5 Many rental stations



First Subtask: A Single Car

For a single car and only two stations, we can efficiently solve all the subproblems of the following form using **Dynamic Programming**:

What is the maximum profit if the car ends at stations s at time t ?

Such a DP-table $DP[s][t]$ has size $2 \times \lceil 100'000/30 \rceil$ and can be filled quickly:


- initialize the cells with $DP[s][0] = 0$ and $DP[s][t] = DP[s][t-1]$
- for all requests r that end at time $30 \cdot t$ at stations s :
 $DP[s][t] = \max(DP[s][t], DP[r.s][r.d/30] + r.p);$
- final answer is $\max(DP[0][MAXT-1], DP[1][MAXT-1])$

This solution does not easily generalize to multiple vehicles.

Second Subtask: Min Cost Max Flow

Instead, let us model it as a flow problem:


Every car is a unit of flow that flows through time and space.



Second Subtask: Min Cost Max Flow

Instead, let us model it as a flow problem:


Every car is a unit of flow that flows through time and space.



Use `cycle_canceling` to handle the negative weights.

Third Subtask: Handle more requests

The usual speedup problem: transform the graph to get non-negative weights so that we can use the faster Min Cost Max Flow algorithm.




Add offset P (= the maximum profit per request) to every timestep.
Note that every s - t -path encounters the same amount of P offsets.

Fourth Subtask: Fine-grained time resolution

If the time resolution is 1 instead of 30 minutes, the graph gets 30 times bigger.


Note: At most 20'000 out of 200'000 vertices have degree > 2 at all.

We can compress these long paths and only create vertices for timesteps where some request starts or ends. This technique is often called *coordinate compression* and reduces the graph size roughly 10 times.



Fifth Subtask: Multiple Stations


Compressing the coordinates per station reduces the graph size again S times.



Implementation detail: Use a `std::map<int,int>` per station to easily map the input times to their compressed coordinates and get the neighboring timestamps.

Fifth Subtask: Multiple Stations

Compressing the coordinates per station reduces the graph size again S times.



Implementation detail: Use a `std::map<int,int>` per station to easily map the input times to their compressed coordinates and get the neighboring timestamps.